

# **Verifiable Computing**

## Part I — Zero Knowledge Proofs

---

Christiano Braga

[cbraga@ic.uff.br](mailto:cbraga@ic.uff.br)

Instituto de Computação  
Universidade Federal Fluminense

**Núcleo de Blockchain**

**Mestrado Profissional do Departamento de Engenharia Elétrica**  
**Universidade de Brasília**

Oct. 1st, 2025

## A useful behaviour in distributed systems

A peer, say  $p_2$ , trusts another peer, say  $p_1$ , in a network application, that a computation was carried out correctly, without  $p_2$  having to redo it or requiring a third-party to certify it.

# Verifiable computing

A *prover* ( $p_1$ ) convinces a *verifier* ( $p_2$ ), in a *protocol*, by means of a **proof**<sup>1</sup>, that a computation was carried out correctly, without the *verifier* redoing it or requiring a third-party to certify it.

---

<sup>1</sup>Let us consider the *intuitive* meaning of the word “proof” for the time being.

## Real-world problems

- **Outsourced computation:** Cloud providers prove that code ran correctly.
- **Sensitive data:** Analyze medical or financial data without exposing it.
- **Blockchains:** Smart contracts require verifiable execution without every node recomputing it.

# Blockchain example

**Problem:** Full on-chain verification is expensive and slow.

**Solution:**

- **Off-chain computation** by a prover.
- **On-chain verification** of proof and public outputs only.

**Benefits:**

- Reduces gas costs dramatically.
- Preserves privacy in smart contracts (when zero-knowledge is used).
- Scales blockchains for complex computations.

# ZK-Rollups

- Bundles (or “rolls up”) thousands of transactions together in an L2.
- Publishes a compressed version back on Ethereum L1.
  - Optimistic rollups assume transactions are valid unless challenged. Examples: [Arbitrum](#) and [Optimism](#)
  - **Zero-Knowledge** rollups, such as [ZKSync](#), use mathematical proofs (called zero-knowledge proofs) to prove that every batch of transactions is valid.

- *Why I support privacy*, blog post dated 2025 Apr 14, Vitalik Buterin, co-founder of Ethereum, concludes:  
[...]  
*supporting **privacy** for everyone, and making the necessary tools open source, universal, reliable and safe is one of the important challenges of our time.*

## Ethereum's future: ZK on RISC-V ii

- *Long-term L1 execution layer proposal: replace the EVM with RISC-V*, dated 2025 Apr 20:

[...]

*In the long term, the primary limiting factors on Ethereum L1 scaling become:*

1. *Stability of data availability sampling and history storage protocols*
2. *Desire to keep block production a competitive market*
3. *ZK-EVM proving capabilities*

*I will argue that replacing the **ZK-EVM with RISC-V** solves a key bottleneck in (2) and (3).*

[...]

- *Simplifying the L1*, blog post dated 2025 May 03, with respect to “Simplifying the consensus layer”, Vitalik says:  
[...]

*The new consensus layer effort (historically called the “beam chain”) aims to use all of our learnings in consensus theory, **ZK-SNARK** development, staking economics and other fields over the last ten years to create a long-term optimal consensus layer for Ethereum.*

[...]

# ZKP: Zero-Knowledge Proofs

You want to prove:

- I know a **secret**  $w$  (called the witness) such that

$$C(x, w) = 0,$$

where

$x$  is the public input,

$w$  is the private witness,

$C$  is the computation you want to prove.

- For example, one could prove one knows the preimage of a hash:  $x$  would be the hash,  $w$  would be the preimage and  $C$  the hash computation.

## Deductive proofs vs Zero-Knowledge Proofs

- A zero-knowledge proof is not a deductive proof, in a mathematical logic sense.
- It is a *probabilistic decentralized certificate* that a computation is successfully carried on.
  - A certificate that a computation is successful, with a tiny tiny chance of error, and that the validity of the certificate can be checked at anytime without a central authority.

## Loan application example i

- Consider you want to apply for a *loan* but do not want to make your portfolio public.
- The computation  $C$  you want to prove is comprised by the *rules* for obtaining a loan.
- The private witness  $w$  is your portfolio.
- The public input  $x$  is your public key, as in an asymmetric key encryption scheme.

## Loan application example ii

- Now, what the algorithms for zero-knowledge proof generation do is, roughly:
  1. Take  $C$ , expressed as *constraint equations*,
  2. generate a *small* proof, *easily verifiable*, that the witness you provided (your portfolio) is a solution to the constraint equations (the loan rules).

## From program to proof: abstract i

- The algorithm we describe below is the so-called Zero-Knowledge Succinct Non-interactive Arguments of Knowledge (zkSNARK), with per-circuit trusted setup.
  - [Groth16](#) is one such algorithm.
- There is a *zoo* of such algorithms, a.k.a proof systems:
  - SNARK, that does not talk about privacy,
  - STARK, that does not require the so-called trusted setup but generates quite large proofs,
  - zkSTARK, STARK with zero-knowledge,
  - and many more.

## From program to proof: abstract ii

1. Express your computation as equations.
2. Transform your equations into a polynomial.
3. Apply cryptographic techniques to make sure a proof for the same computation can not be forged.
4. Constructs a proof.
  - For zkSNARK, it is as simple as 3 points in a curve.
5. The verifier checks the proof, using the public inputs.

## \*From program to proof: “concrete” i<sup>2</sup>

Of course, there are still many details abstracted away.

1. Express the computation in the so-called Rank 1 Constraint System (**R1CS**) notation.
- Turn your program into a set of Rank-1 Constraint System (R1CS) equations:

$$(A_i \cdot s) \cdot (B_i \cdot s) = (C_i \cdot s)$$

for each constraint  $i$ , where  $s$  is a vector of all variables (inputs, witness, intermediates).

- It shows that secret data  $w$  satisfies the constraints describing the computation.

---

<sup>2</sup>This, and the next slides tagged with \* before their headings, may be skipped if you do not want more details about proof generation.

## \*From program to proof: “concrete” ii

2. Convert to a Quadratic Arithmetic Program (**QAP**).
  - Transform the R1CS into polynomials:
    - Define polynomials  $A(X)$ ,  $B(X)$ ,  $C(X)$  that encode the constraints.
    - The prover must show that for some  $s$ ,  $(A(s) \cdot B(s) - C(s))$  is divisible by  $Z(s)$  (vanishing polynomial) where  $Z(X)$  is the target polynomial whose roots are the constraint indices.
  - This becomes an algebraic check: proving you know  $s$  making the constraints hold.
    - This is essential for a fast verification of the proof.

### 3. Trusted setup (Structured Reference String, SRS)

- The system creates cryptographic parameters.
  - Some of these parameters are called *toxic waste*.
  - *If the toxic waste is leaked, an attacker could forge proofs.*
- This setup is done in two steps: a general one and once per circuit, that is, for each  $C$ ,  $w$  and  $x$ .

## \*From program to proof: “concrete” iv

### 4. Prover creates the proof

- The prover:
  - Computes a commitment (a hash, essentially) to the witness and polynomials using the SRS.
  - Constructs a very small proof. (In SNARK, **three elliptic curve points**:  $(\pi_A, \pi_B, \pi_C)$ .)
- These points encode the check that:
  - The prover knows  $w$  such that  $C(x, w) = 0$ .
  - The witness satisfies the **polynomial divisibility** condition.
- Thanks to SRS, the proof stays tiny and fast to verify.

## \*From program to proof: “concrete” v

### 5. Verifier checks the proof

- The verifier uses:
  - The public input  $x$ ,
  - The public SRS (setup parameters),
  - The succinct proof  $(\pi_A, \pi_B, \pi_C)$ .
- And performs a small number of **pairing checks on elliptic curves**.
- Pairings let the verifier efficiently check polynomial relationships encoded in the proof.
- If the checks pass, the verifier is convinced:
  - The prover knows a witness  $w$ ,
  - The statement  $C(x, w) = 0$  is true,
  - without ever learning  $w$ ...

## \*Why a polynomial?

- They provide a proper mathematical foundation, together with the domain of prime finite fields, to represent a computation.
- Are amenable to the many cryptographic techniques used in the proof generation and its verification.

## \*Why do we express the computation as arithmetic constraints over a prime field?

- In ZKP, we want to prove “I know a witness that satisfies some statement,” in a way that is Succinct (small proof), efficient to verify, and zero-knowledge (doesn’t leak the witness).
- To do that with zk-SNARKs one needs to:
  - Go from arithmetic circuits  $\rightarrow$  R1CS  $\rightarrow$  QAP  $\rightarrow$  polynomial commitments.
  - All of these fundamentally need arithmetic over a finite field.

## \*And here's why i

1. Need a closed algebraic structure for constraints:
  - We want to represent SAT problems with arithmetic circuits as they are less verbose.
  - Arithmetic circuits compute with addition, multiplication, multiplication by constants.
  - We want:
  - For every pair of elements in the carrier set of the algebra, addition and multiplication are defined.
  - Every non-zero element has a multiplicative inverse.
  - A finite field  $\mathbb{F}_p$ , with  $p$  a prime, gives exactly this:
    - No undefined operations (like division by 0 in inverses, except on 0 itself).
    - Closed, consistent, total arithmetic.
2. Constraints as polynomial equations that vanishes in points that satisfy all constraints simultaneously.

## \*And here's why ii

- In R1CS / QAP, the statement “the computation was done correctly” becomes: A set of polynomial equations over the variables and inputs  $(A_i \cdot s) \cdot (B_i \cdot s) = C_i \cdot s$ .
- This only makes mathematical sense in a ring where addition, multiplication, and inverses are well defined.
- We can do polynomial interpolation and division.

**Finite fields fit perfectly.**

# A circuit for loan eligibility i

## Arithmetic constraint:

$$\text{salary} \geq \text{salary\_threshold} \wedge \text{credit\_score} \geq \text{credit\_threshold}$$

- Public thresholds allow for circuit reuse.
- Private inputs remain hidden.

## A circuit for loan eligibility ii

### Inputs:

- Private: salary, credit\_score
- Public: salary\_threshold, credit\_threshold

### Computation:

$$is\_salary\_ok = salary \geq salary\_threshold$$

$$is\_credit\_ok = credit\_score \geq credit\_threshold$$

$$loan\_approved = is\_salary\_ok \times is\_credit\_ok$$

## A circuit for loan eligibility iii

### R1CS:

- Let  $x_1 = \text{is\_salary\_ok}$  and  $x_2 = \text{is\_credit\_score\_ok}$

```
x1 * (x1 - 1) = 0           # boolean constraint  
x2 * (x2 - 1) = 0  
loan_approved - x1 * x2 = 0 # AND operation
```

# A circuit for loan eligibility iv

## R1CS: (cont.)

- However, R1CS does not know how to represent inequalities.
  - We need to transform the quantities into bitstreams of some size and subtract them.
  - The code below implements “less than” operation ( $>$ ) in [Circom](#), the first domain-specific language (DSL) for ZK programming.

```
template LessThan(n) {
    assert(n <= 252);
    signal input in[2];
    signal output out;
    component n2b = Num2Bits(n+1);
    n2b.in <== in[0]+ (1 << n) - in[1];
    out <== 1-n2b.out[n];
}
```

## ZK virtual vachines

- Execute arbitrary programs inside a ZK proof system.
- Example: [Risc0 zkVM](#), a zkSTARK VM that emulates RISC-V.
- From Risc0 [helloworld](#) example:
  - zkVM applications are organized into a host program and a guest program.
  - The host first executes the guest program and then proves the execution to construct a receipt.
  - The receipt can be passed to a third party, who can examine the journal to check the program's outputs and can verify the receipt to ensure the integrity of the guest program's execution.

# ZK Domain-specific languages

- **Purpose:** Simplify writing circuits targeting ZK proof systems.
- **Examples:**
  - [Circom](#): the first DSL for ZK programming.
  - [Clean](#): a DSL implemented on top of Lean programming language.
  - [Noir](#): circuits are written in Rust.
  - [Lurk](#): circuits are written in a Lisp dialect.

## My own adventures with ZKP

- A simple credit-score proof generation and verification in Lurk.
- A prototype Metamath checker in Lurk.
- **ZKForAll**: Proving soundness and completeness of ZK circuits in Clean.

## Summary

- Verifiable computing is essential for secure, private, and scalable decentralized systems.
- Verifiable computing allows trustless verification of computations without re-execution.
  - **ZKP** is one such technique.
  - Prove computation correctness without revealing inputs.
- Arithmetization and R1CS bridge traditional computation and ZK proofs.
- SNARKs and STARKs provide succinct proofs; Groth16 is a widely used zkSNARK.
- zkVMs like Risc0 enable general-purpose verifiable execution.
- zk DSLs (Noir, Lurk, Clean) simplify writing provable programs.

# Conclusion

- This was just a very small appetizer to the world of verifiable computing with ZKP.
- A lot of work to be done!
- ZKP brings a universe of possibilities to computing, in distributed systems in particular.
- Join the ZK revolution: privacy + scalability = the future of computing!

## Acknowledgements

A big thank you to [Semar Augusto Martins](#) and [Arthur Paulino](#) for their various comments on an earlier version of these slides.